

CSC263H

Data Structures and Analysis

Prof. Bahar Aameri & Prof. Marsha Chechik

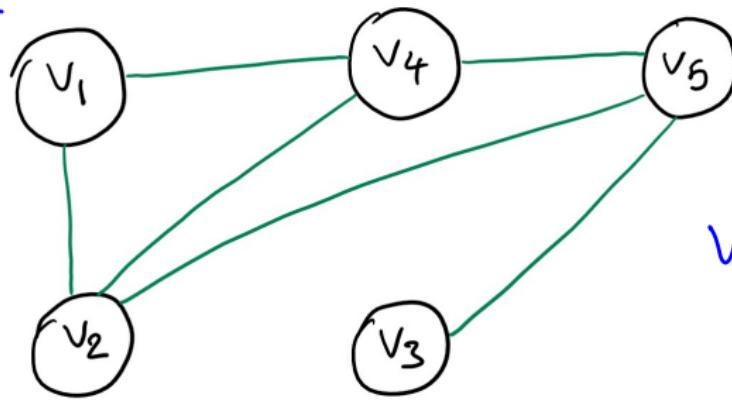
Winter 2024 – Week 8

Graph ADT

A **graph** is a tuple of two sets $G = \langle V, E \rangle$, where V is a set of **vertices**, and E is a set of **edges** between those vertices.

E itself is a set of tuples (v_i, v_j) , where $v_i, v_j \in V$.

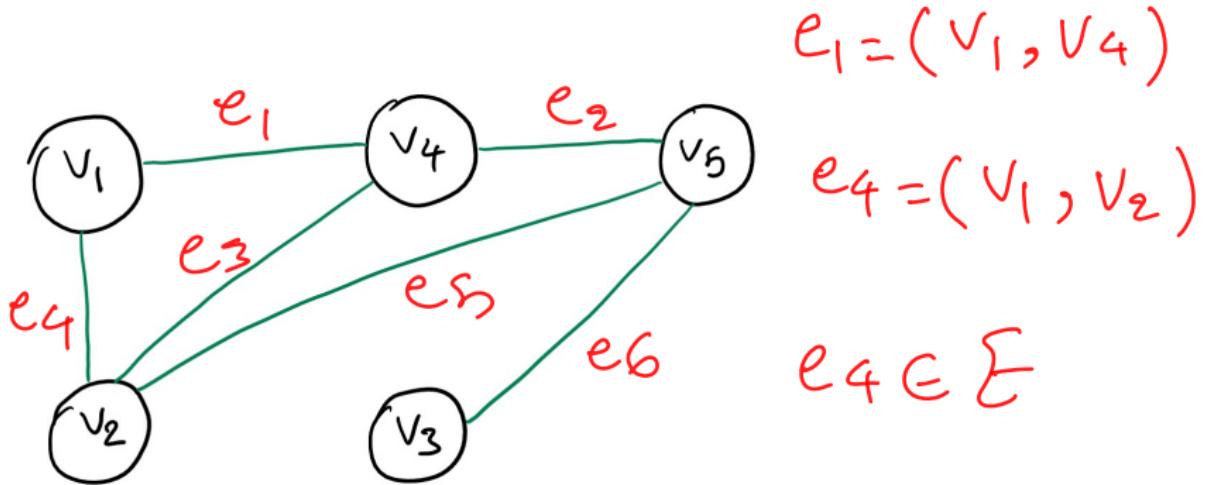
G



$$G = \langle V, E \rangle$$

$$V = \{v_1, v_2, v_3, v_4, v_5\}$$

$$E = \{(v_1, v_2), (v_1, v_4), (v_1, v_5), (v_2, v_3), (v_2, v_4), (v_2, v_5), (v_3, v_5)\}$$



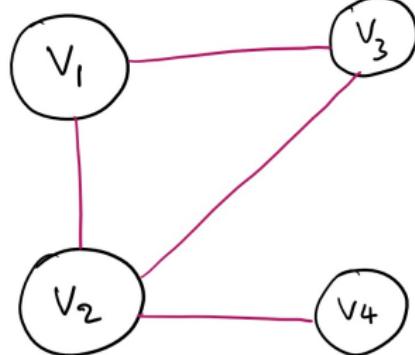
- Network Traffic flow
- Social Webs
- Task Scheduling (Operating Systems)
- Maps and GPS
- VLSI and Chip Design

Graph ADT: Basic Definitions

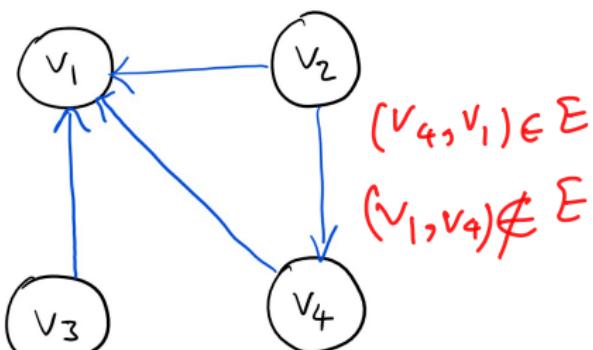
A **directed graph** is a graph in which edges have orientations.

That is, the order of tuples matter: (v_i, v_j) and (v_j, v_i) represent **different** edges.

An **undirected graph** is a graph in which edges have no orientation. That is, the order of tuples does **not** matter: (v_i, v_j) is the same as (v_j, v_i) .



$$(v_i, v_j) = (v_j, v_i)$$



$$(v_i, v_j) \neq (v_j, v_i)$$

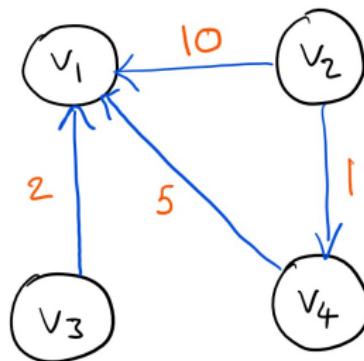
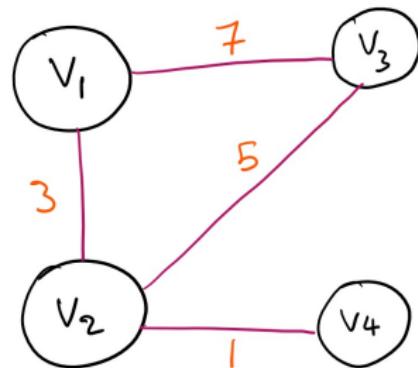
- Facebook friendship can be represented by undirected graphs:
 - Each person (account) is represented by a vertex;
 - If u and v are friends iff there exists an edge between u and v .
- Instagram followers can be represented by directed graphs.
 - Each person (account) is represented by a vertex;
 - If u follows v iff there exists an edge from u to v .

Note: Usually, if neither type of graph is specified, our default will be an undirected graph.

Graph ADT: Basic Definitions

A **weighted graph** is a graph in which a weight (which is a number) is assigned to each edge.

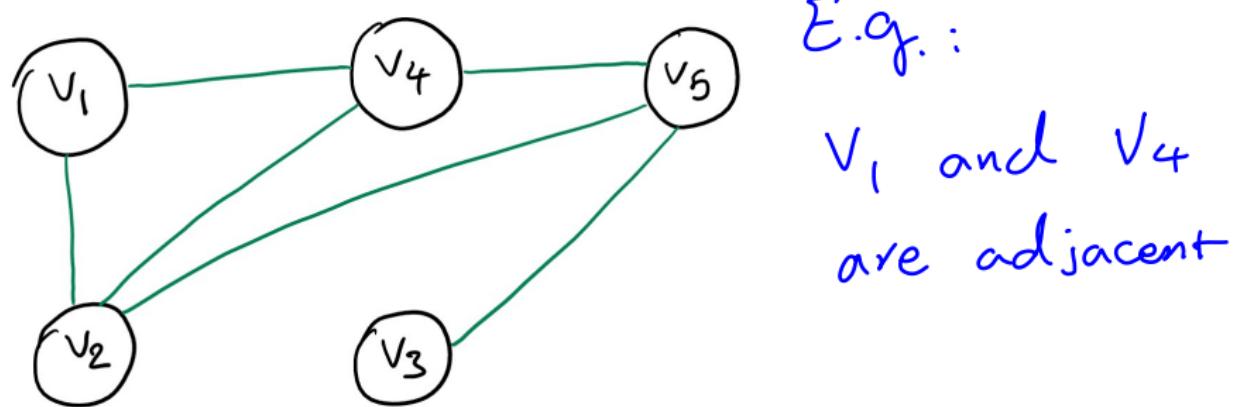
Note: Weights of edges might represent costs, lengths or capacities, depending on the problem at hand.



- Network traffic flow can be represented by weighted directed graphs:
 - Each hub is represented by a vertex;
 - Incoming traffic from a hub u to a hub v is represented by an incoming edge from u to v ;
 - Outgoing traffic from a hub v to a hub z is represented by an outgoing edge from v to z ;
 - Traffic flow is represented by the weight of each edge.
- Maps can be represented by weighted graphs:
 - Each address point is represented by a vertex;
 - Roads between address points are represented by edges;
 - The distance between two points is represented by the weight of the edge between the points.

Note: Our default is that a graph is NOT weighted, unless otherwise is specified.

Two vertices are **adjacent** if there is an edge between them.
In this case, we also say that the two vertices are **neighbours**.

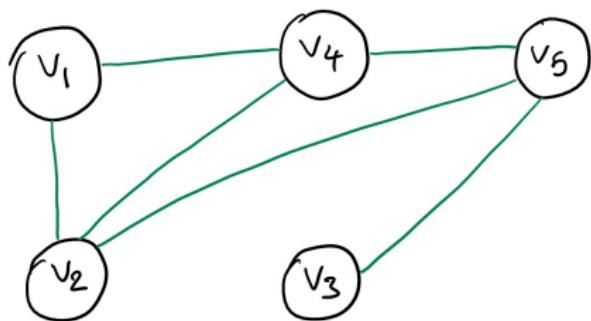


Graph ADT: Basic Definitions

A path between vertices u and w is a sequence of distinct edges $(v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k)$, where $u = v_0$ and $w = v_k$, and for all tuples (v_i, v_j) in the sequence, v_i and v_j are distinct.

The length of the path is the number of edges in the path.

Start at vertex u , take an edge to get to a new vertex v_1 , take another edge to get to another new vertex v_2 , and keep following edges until arriving at v_k .

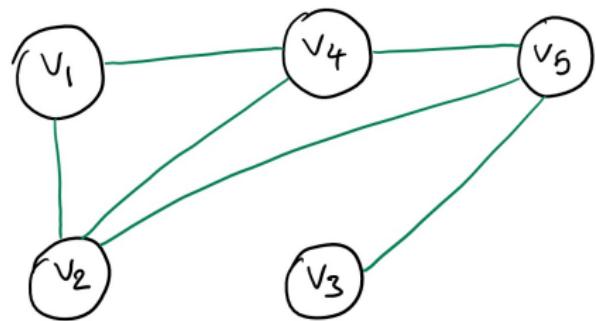


Path 1: (v_1, v_4)
length: 1

Path 2: $(v_1, v_2)(v_2, v_4)$
length: 2

Path 3: $(v_1, v_2)(v_2, v_5)(v_5, v_4)$
length: 3

The **distance** between two vertices is the length of the **shortest path** between the two vertices.
The distance between a **vertex and itself** is always 0.



distance of v_1 and v_4 : 1

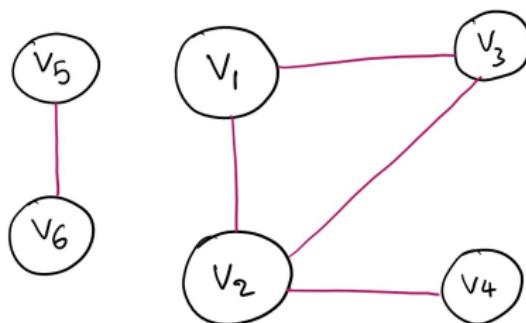
" " v_1 and v_5 : 2

" " v_1 and v_3 : 3

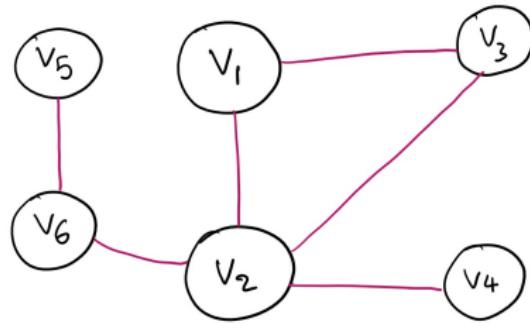
Graph ADT: Basic Definitions

A graph is **connected** if there is a **path** between **every pair** of vertices in the graph. Otherwise, the graph is **disconnected**.

G_1 X

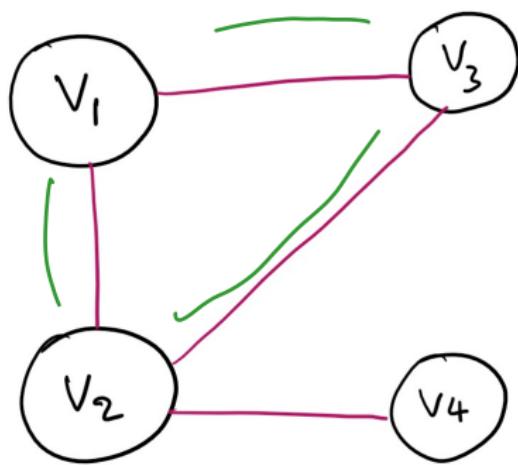


G_2 ✓



Minimum number of edges in an undirected **connected** graph with n vertices: $n - 1$

In a graph, a **cycle** is path from a vertex to itself.

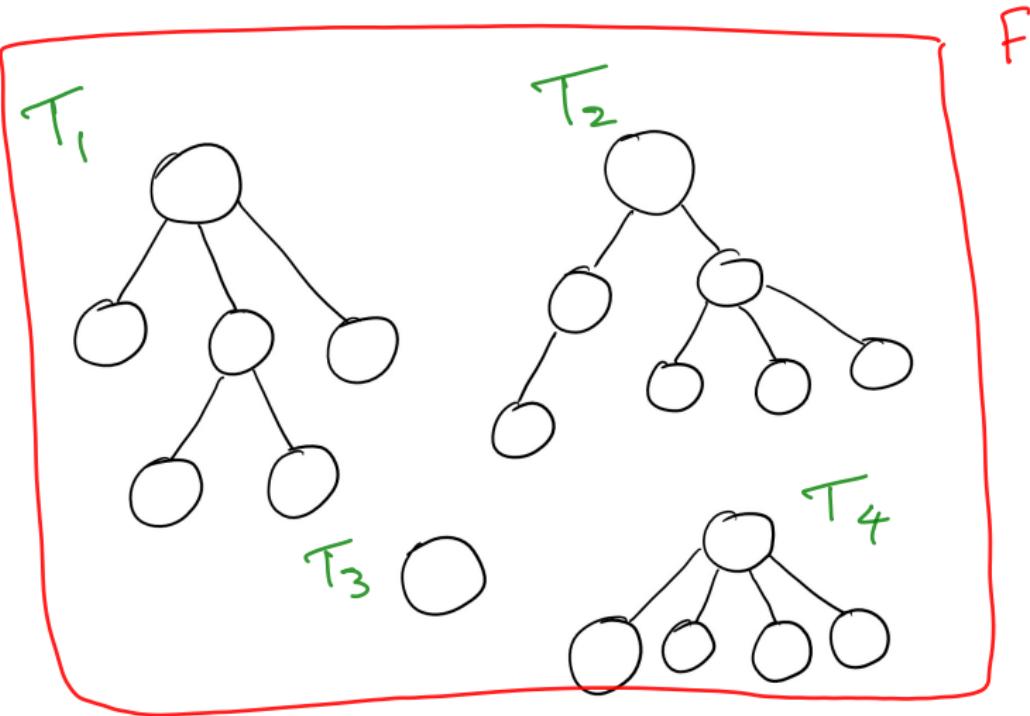


ASSumption:
Length of a cycle
is at least 3

Graph ADT: Basic Definitions

A **tree** is **connected** graph and contains **no cycle**.

A **forest** is a **collection of disjoint trees**.

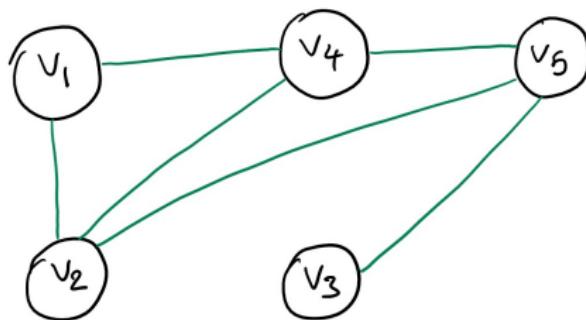


Data Structures for Graphs: Adjacency Matrix

Let $V = \{v_1, v_2, \dots, v_n\}$ ($|V| = n$).

Let A be a $n \times n$ matrix.

$$A[i, j] = \begin{cases} 1 & \text{if } (v_i, v_j) \in E \\ 0 & \text{otherwise} \end{cases}$$



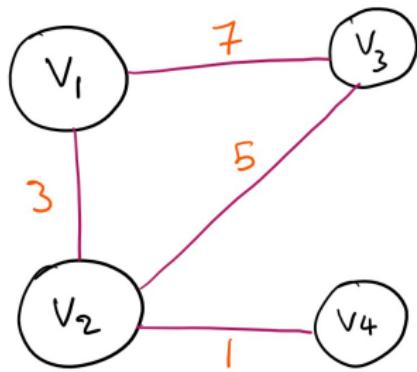
	1	2	3	4	5
1	0	1	0	1	0
2	1	0	0	1	1
3	0	0	0	0	1
4	1	1	0	0	1
5	0	1	1	1	0

For an undirected graph, adjacency matrix is **symmetric**: $A[i, j] = A[j, i]$.

For a **weighted** graph:

If $(v_i, v_j) \in E$, store $w(v_i, v_j)$ in $A[i, j]$ ($w(v_i, v_j)$ denotes the weight of the edge (v_i, v_j));

If $(v_i, v_j) \notin E$, store $-1/0/\infty$ (depending on application).



1	2	3	4
2	3	7	-1
3	5	-1	1
4	-1	1	-1

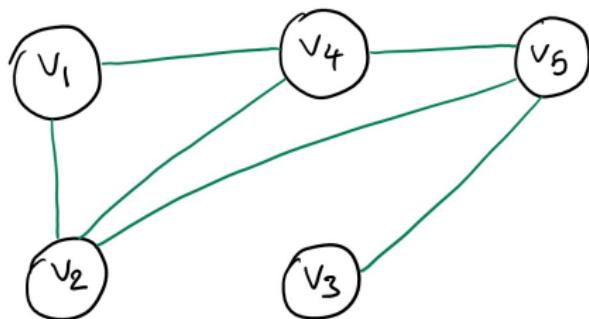
Data Structures for Graphs: Adjacency Lists

Let $V = \{v_1, v_2, \dots, v_n\}$ ($|V| = n$).

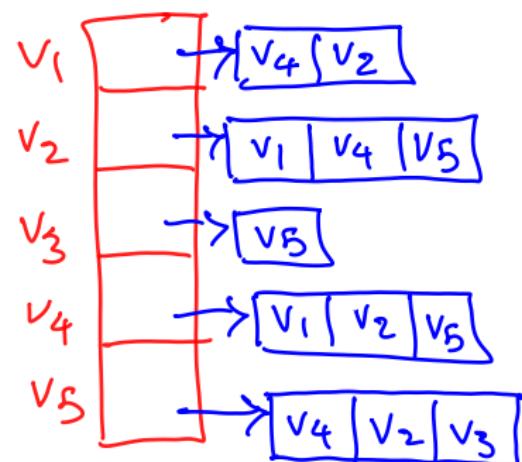
Let L be a list of size n .

Each position $L[i]$ corresponds with a vertex v_i and stores a list A_i of all vertices that have an edge from v_i :

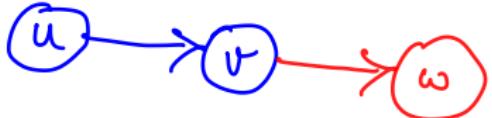
A_i includes v_j iff $(v_i, v_j) \in E$



For an undirected graph, edge (u, v) is stored **twice** (u in sub-list corresponding with v and v in sub-list corresponding with u).



- Add/Remove a vertex
- Add/Remove an edge.
- Edge Query: given vertices u, v , does G contain edge (u, v) or (v, u) ?
- Neighbourhood (undirected graph): given vertex v , return set of vertices u such that $(v, u) \in E$.
- In-neighbourhood / out-neighbourhood (directed graph): given vertex v , return set of vertices u such that $(u, v)/(v, u) \in E$.
- Degree/ In-degree/ Out-degree:
size of neighbourhood / in-neighbourhood / out-neighbourhood.
- Traversal: visit each vertex of a graph to perform some task.



Adjacency Matrix vs Adjacency Lists

$$|V|=n \quad |E|=m$$

	Adjacency Matrix	Adjacency Lists
Space Complexity	$\Theta(n^2)$	$\Theta(n+m)$
Add/Remove a Vertex	$\Theta(n)$	$\Theta(1) \rightarrow$ $\Theta(m) \rightarrow$ add remove
Edge Query	$\Theta(1)$	$\Theta(n)$ more precise bound:
Neighbourhood	$\Theta(n)$	$\Theta(n)$ \rightarrow $\Theta(\min(n, m))$

For an arbitrary undirected graph

Choose the more appropriate implementation depending on the problem.

Graph Traversal:

Visit **all** the vertices in a graph to perform some task (task depends on the application).
The visits can follow **different orders**.

Graph Traversal Algorithms: Starting from a given vertex $s \in V$, visit every vertex in the graph that is reachable from s .

- **Breadth-First Search (BFS):**

- **Intuition:** Starting from s , first visit **all** (unvisited) neighbours of s , then visit **all** (unvisited) neighbours of all neighbours of s , then visit **all** (unvisited) neighbours of all neighbours of all neighbours of s , and so on, until all reachable vertices are visited.

- **Depth-First Search (DFS):**

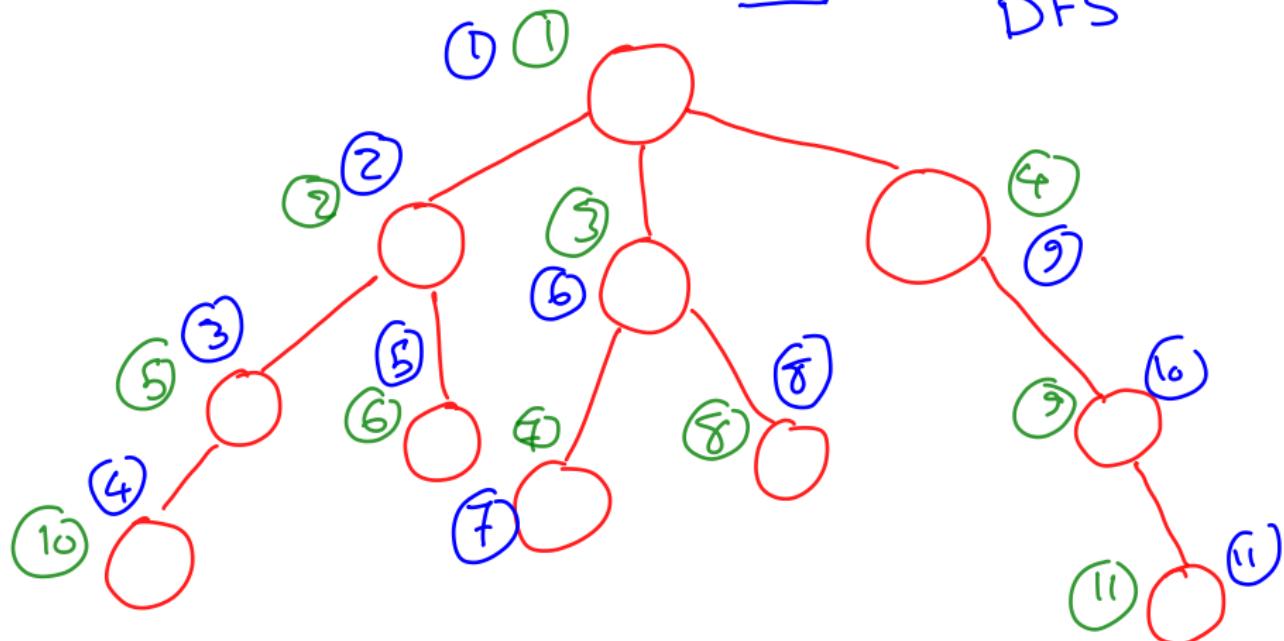
- **Intuition:** Traverse the depth of any particular path before exploring its breadth. Starting from s , first visit **a neighbour** n_1 of s , then visit **an unvisited neighbour** n_2 of n_1 , then visit **an unvisited neighbour** n_3 of n_2 , and so on. If visiting a neighbour n_k of n_{k-1} and all neighbours of n_k are already visited, try another neighbour of n_{k-1} . Stop when all vertices reachable from s are visited.

BFS vs DFS: A Familiar Example

Consider performing the BFS and DFS algorithms on the root of a tree.

S

BFS
DFS

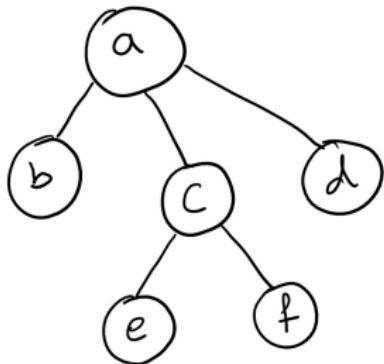


BFS in a tree (starting from root) is a **level-by-level** traversal.

DFS visits the **child vertices** before visiting the sibling vertices.

NotYetBFS($T, root$):

- 1 $Q = \emptyset$
- 2 Enqueue($Q, root$)
- 3 **While** Q not empty:
- 4 $u = \text{Dequeue}(Q)$
- 5 **print** u
- 6 **for** each child c of u :
- 7 Enqueue(Q, v)

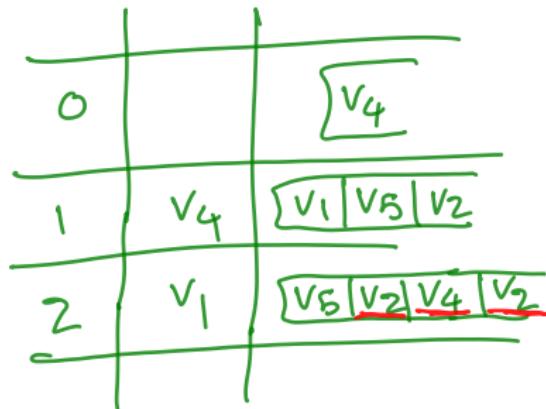
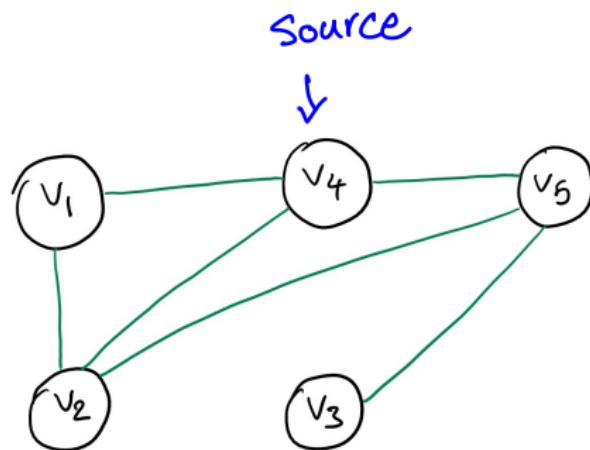


Iteration	Output	Queue
0	-	a
1	a	b c d
2	b	c d
3	c	d e f
4	d	e f
5	e	f
6	f	

BFS Implementation

What happens if we run NotYetBFS on a non-tree graph?

BFS will try to visit some of the vertices **twice**.



Note: We refer to the **starting vertex** of a BFS call as the **source** vertex.

How avoid visiting a vertex **twice**?

Remember the visited vertices by **labelling** them using colours.

- **White:** **Unvisited** (undiscovered) vertices.
Have **not** been enqueued.
- **Gray:** **Encountered** (discovered) vertices.
Have been **enqueued**.
- **Black:** **Explored** vertices.
Have been **dequeued** and all of their **neighbours** are **enqueued**.

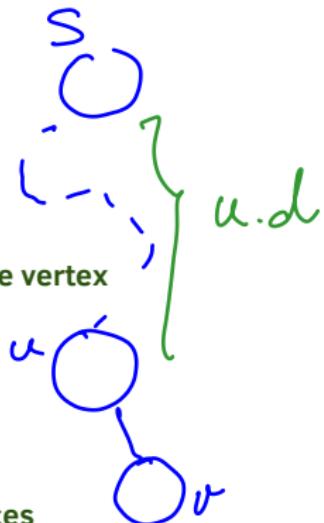
- **Initially** all vertices are **white**.
- Change a vertex's color to **gray** the first time **visiting** (enqueueing) it.
- Change a vertex's color to **black** when all its **neighbours** have been **encountered** (enqueued).
- **Avoid** visiting (enqueueing) **gray** or **black** vertices.
- In the **end**, all vertices that are reachable from the source are **black**.

Other useful values to remember during the traversal:

- The vertex from which v is encountered. (stored in $v.p$)
- The **distance value**: the distance **from v** to the **source** vertex of the BFS. (stored in $v.d$)

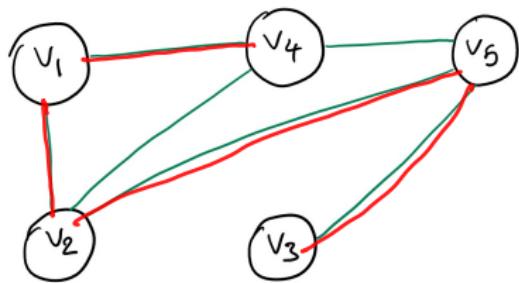
BFS(G, s):

1. **for each vertex $v \in V - \{s\}$:** # Initialize vertices
2. $v.colour = \text{White}$
3. $v.d = \infty$
4. $v.p = \text{Nil}$
5. $Q = \emptyset$
6. $s.colour = \text{Gray}$ # start BFS by encountering the source vertex
7. $s.d = 0$ # distance from s to s is 0
8. Enqueue(Q, s)
9. **While Q not empty:**
10. $u = \text{Dequeue}(Q)$
11. **for each $v \in G.\text{adj}[u]$:**
12. **if $v.colour == \text{White}$** # only visit unvisited vertices
13. $v.colour = \text{Gray}$
14. $v.d = u.d + 1$ # v is "1-level" farther from s than u
15. $v.p = u$ # v is introduced as u 's neighbour
16. Enqueue(Q, v)
17. $u.colour = \text{Black}$ # u is explored as all its neighbours have been encountered



The blue lines are the same as NotYetBFS.

Source : v_1



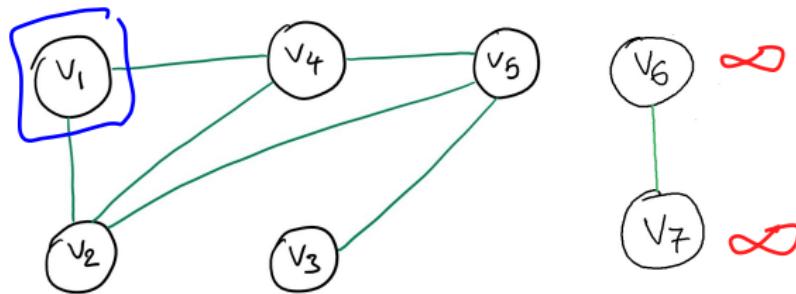
Vertex	colour	d	p
v_1	GB	0	Nil
v_2	WGB	1	Nil
v_3	WGB	3	Nil
v_4	WGB	1	Nil
v_5	WB	2	v_2

Iteration	Queue	u
0	v_1	
1	$v_2 v_4$	v_1
2	$v_4 v_5$	v_2
3	v_5	v_4
4	v_3	v_5
5		v_3
6		

The set of all edges between vertices and their **introducer** form a tree, call a **BFS tree**

What if G is disconnected?

G



The **infinite distance** values of v_6 and v_7 (i.e., $v_6.d$ and $v_7.d$) indicate that they are unreachable from the source.

BFS Worst-case Running Time

The total amount of work (use adjacency list):

$$n = |V| \quad m = |E|$$

- When visiting each vertex: n

Enqueue, Dequeue, assign values to $v.colour$, $v.d$, $v.p$, etc. $\Theta(1)$

- At each vertex, check all its neighbours (i.e., all its incident edges).

Each edge is checked at most twice (by the two end vertices)

- Total for 1: $\Theta(n)$

- Total for 2: $\Theta(m)$

Total running time: $\Theta(n + m)$

Exercise: What is the BFS Worst-case running time when using an adjacency matrix? $\Theta(n^2)$

BFS Properties

- BFS can be performed on both **directed and undirected** graphs.
- What do we get after performing a BFS:
 1. We get to visit every vertex which is **reachable** from the **source** s .
 2. We generate a **BFS tree**.

Each edge in a BFS tree is called a **BFS tree edge**.
The BFS tree connects **all** vertices, if the graph is **connected**.
That is, if the graph is **not** connected, then for **every** BFS tree of the graph, **some** vertices of the graph are **not included** in the tree.
 3. For every vertex v , $v.d$ stores the minimum distance (or **shortest-path** distance) between the BFS **source** and v .

To get the **shortest-path** itself, follow $.p$ attribute (i.e., the red edges in the example) from v to s backward.
(Proof: Theorem 22.5 in CLRS)



- **Traversing** graphs (Based on Property 1).
- Identifying **connectedness** (Based on Property 2).
- Finding **shortest-path** between two points (Based on Property 3).

- Relevant sections and exercises from Chapter 6 of the course notes.
- After-lecture Readings: CLRS Sections 22.1, 22.2