

CSC263H

Data Structures and Analysis

Prof. Bahar Aameri & Prof. Marsha Chechik

Winter 2024 – Week 4

- For a binary tree with n nodes, what is the smallest possible height?

$$\Theta(\lg n)$$

- What kinds of binary trees have such height?

Complete Binary Trees

Data Structures for Dictionaries: AVL Trees

In a **complete** binary tree, the **heights** of the **left** and **right** sub-trees of any node differ by at most 1.

Balance Factor (BF): The height of the right sub-tree minus the height of the left sub-tree.

$$BF(n) = n.right.height - n.left.height$$

AVL Invariant: A node n satisfies the AVL invariant if $-1 \leq BF(n) \leq 1$.

AVL-Balanced: A binary tree that all of its nodes satisfy the *AVL invariant*.



AVL Tree: A **BST** which is **AVL-balanced**.

Invented by Georgy Adelson-Velsky and E. M. Landis in 1962.



Implication: In an AVL tree, the BF of every node is -1, 0, or 1.

(Note: Height is measured by the number of levels (number of nodes in the longest path from the root to a leaf).)

Data Structures for Dictionaries: AVL Trees



Data Structures for Dictionaries: AVL Trees



- If $BF(x) = +1$, x is **right heavy**.
- If $BF(x) = -1$, x is **left heavy**.
- If $BF(x) = 0$, x is **balanced**.

Theorem: The height of an AVL tree with n nodes is at most $1.44 \log_2 (n + 2)$.

\Rightarrow For an AVL tree with height h we have: $h \in \Theta(\log n)$.

Storage:

In addition to $x.key$, $x.left$, $x.right$, $x.p$, $x.height$ is stored in each node x .

Operation Implementation:

- *AVLSearch*: Same as *BSTSearch*.

- *AVLInsert* and *AVLDelete*:


Challeng: Keeping tree *balanced* after each update (insert/delete).

1. Maintain AVL invariant for all affected nodes (i.e., ancestors of the inserted/deleted node).
2. Maintain the BST property.
3. Update height of the affected nodes accordingly.

AVL Trees – Implementation


Insert 35

Insert 20



Observations:

1. Inserting/deleting a node can only change the balance factors of its ancestors.
2. Inserting/deleting a node can cause a *sub-tree's height* to increase/decrease by at most 1.
So the balance factor of the affected nodes changes by at most 1.
The balance factor of the affected nodes can only be -2 or 2.




Assumption: y is the lowest ancestor that became unbalanced.
That is, all decedents of y satisfy the AVL invariant.

Case 1: All nodes (in the subtree) except y satisfy the AVL invariant and $BF(y) = 2$

- To rebalance, must **increase** the height of the **left subtree** of y and **decrease** the height of the **right subtree** of y .

Case 2: All nodes (in the subtree) except y satisfy the AVL invariant and $BF(y) = -2$.

(symmetric to Case 1) *Assumption: x is either balanced or right-heavy*




Assumption: y is the lowest ancestor that became unbalanced.
That is, all decedents of y satisfy the AVL invariant.

Case 1: All nodes (in the subtree) except y satisfy the AVL invariant and $BF(y) = 2$.

Case 2: All nodes (in the subtree) except y satisfy the AVL invariant and $BF(y) = -2$.
(symmetric to Case 1)


- To rebalance, must **increase** the height of the **right subtree** of y and **decrease** the height of the **left subtree** of y .

$h_2 - h_1 = -2$ Assumption: x is either balanced or left-heavy



\Rightarrow

$h_1 - 1$



$$BF(x) = h_{2+1} - h_{1-1} = -2 + 2 = 0$$

AVL Rebalancing: Right Rotation



Rebalancing Move: **Rotation!**

Requirements:

1. Changes heights of a node's left and right subtrees.
2. Maintains BST property.

Right rotation around y

BST order to be maintained: $A x B y C$



AVL Rebalancing: Left Rotation



Rebalancing Move: **Rotation!**

Requirements:

1. Changes heights of a node's left and right subtrees.
2. Maintains BST property.

Left rotation around γ


BST order to be maintained: $C \gamma B x A$




AVL Insertion: Case 1

Case 1: Inserting a new node into the right-subtree of y while y is right-heavy (i.e., $BF(y) = +1$) before insertion.


- **Case 1.1:** Insert the new node to the right subtree of x (x is the right child of y)
- **Case 1.2:** Insert the new node to the left subtree of x



AVL Insertion: Case 1.1




AVL Insertion: Case 1.2



Rotation works for adjusting the heights of the *left side* and the *right side*. But height of the middle subtree does *NOT* shrink when rotating around root y .



Must move the new node to the side first!

AVL Insertion: Case 1.2



Both cases can be fixed with a double
right-left rotation


AVL Insertion: Case 1.2



AVL Insertion: Case 1.2

Left rotation
around y

\Rightarrow



$$BF(y) = 2$$

Case 1:

- **Case 1.1:** single left rotation.
- **Case 1.2:** double right-left rotation.

Case 2: (symmetric to Case 1)


- **Case 2.1:** single right rotation.
- **Case 2.2:** double left-right rotation.

→ if x is right-heavy


→ if x is left-heavy

→ if x is left-heavy

→ if x is right-heavy




$$BF(y) = -2$$




AVL Insertion: Example

Rotation




right-rotation
around 2

\Rightarrow



\Rightarrow left-rotation
around γ



Operation Implementation:

- *AVLSearch*: Same as *BSTSearch*.

- *AVLInsert* and *AVLDelete*:

Challeng: Keeping tree *balanced* after each update (insert/delete).

1. Maintain AVL invariant for all affected nodes (i.e., ancestors of the inserted/deleted node).
2. Maintain the BST property.
3. Update height of the affected nodes accordingly:
Update heights going up from the **new leaf** to the **root**.

Observations:


1. *BSTInsert/BSTDelete* already traverse exactly the nodes which are ancestors of the modified node.
2. If ^D is an AVL tree before the recursive call, then so are the trees rooted at *D.left* and *D.right*.
3. If we can ensure that after the recursive call the trees rooted at *D.left* and *D.right* are still AVL trees, then we can apply rotations to fix the tree rooted at *D*.

AVL Insert: Implementation

Implementation Idea:

1. *AVLInsert*(*root*, *x*):
Precond: The tree rooted at *root* is an AVL tree.
Postcond: The tree rooted at *root* is an AVL tree that includes node *x*.
2. After the recursive call, all the decedents of *root* satisfy the AVL invariant. All we need to do is to fix the tree rooted at *root* by applying rotations.

```
AVLInsert(root, x):  
1   if root == NIL :  
2       root = x  
4   else if root.key > x.key :  
5       AVLInsert(root.left, x)  
6   else :  
7       AVLInsert(root.right, x)  
8   BF = root.right.height - root.left.height  
9   if BF < -1 or BF > 1 :  
10      # Fix the imbalance for the root node  
11      fix_imbalance(root)  
12  root.height = max(root.right.height, root.left.height) + 1
```



AVL Insert: Outline

1. Insert like a BST.

$$\rightarrow \Theta(h) = \Theta(\log n)$$

2. If still balanced, return.

3. Else: (need re-balancing)

• **Case 1:**

left

– **Case 1.1:** single ~~right~~ rotation.

– **Case 1.2:** double ~~left-right~~ rotation.

right-left

• **Case 2:** (symmetric to Case 1)

right

– **Case 2.1:** single ~~left~~ rotation.

– **Case 2.2:** double ~~right-left~~ rotation.

left-right

4. **Updated** the height of affected nodes.

$$\Theta(1)$$


$$\Theta(\log n)$$

Worst-case running time:


$$\Theta(\log n)$$

AVL Insert: Example


Insert 20 and 27 into the following AVL tree



left rotation
around 25
⇒



right rotation
around 40



AVL Delete: Outline

1. Delete like a BST.

$\rightarrow \Theta(\lg n)$

2. If still balanced, return.

3. Else: (need re-balancing)

• **Case 1:**

– **Case 1.1:** single left rotation.

– **Case 1.2:** double right-left rotation.

• **Case 2:** (symmetric to Case 1)

– **Case 2.1:** single right rotation.

– **Case 2.2:** double left-right rotation.


4. **Updated** the ~~balance factors~~ of affected nodes.

height


$\Theta(\lg n)$
Worst-case running time:

$\Theta(\lg n)$

Delete 14



Left rotation
around 17



- We **augmented** BSTs by storing additional information (the height) at each node.
 - The additional information enabled us to keep the tree balanced.
 - We could maintain this additional information efficiently in modifying operations (i.e., without affecting the running time of *Insert* or *Delete*).

Augmented Data Structure: A **modification** of an **existing** data structure by storing additional information and/or performing additional operations.

- Why Augmentation is needed?
 - Textbook data structures rarely satisfy what is needed for solving real problems.
 - It is rarely needed to invent something completely new.
 - Augmenting known data structures to serve specific needs is the sensible middle-ground.

General procedure:

1. Choose data structure to augment.
2. Determine additional information.
3. Check additional information can be maintained during each original operation (and additional cost, if any).
4. Implement new operations.


Ordered Sets:

1. *Insert, Delete, Search*: Same as Dictionaries.
2. *Rank(k)*: return *rank* of key k , i.e., index of k in sorted ordering of set elements.
3. *Select(r)*: return key with rank r .

Example: For the set $\{27, 56, 30, 13, 15\}$,

$\text{Rank}(15) = 2$ and $\text{Select}(4) = 30$

because the sorted order is $[13, 15, 27, 30, 56]$.



Tutorial 4: Implementation of Ordered Sets by **Augmented** AVL trees.

Theorem (AVL tree augmentation): In augmenting AVL trees, if the additional information of a node only depends on the information stored in its children and itself, this information can be maintained efficiently during *AVLInsert* and *AVLDelete* without affecting their $\Theta(\log n)$ worst-case runtime. (Proof Similar to Theorem 14.1 of CLRS)

- After-lecture Readings: Notes on AVL trees (posted on portal).
- Review AVL trees in the Course Notes (Chapter 3).
- Example Exercises (Course notes).
- Detailed implementation of *AVLDelete*.